

Edge-Based Cloud Computing as a Feasible Network Paradigm

Joe Elizondo and Samuel Palmer

Department of Computer Science University of Texas at Austin

Email: {elizondo, spalmer}@cs.utexas.edu

December 3, 2009

Abstract

The term edge based cloud computing refers to a network of edge systems that provide the services currently provided by data center clouds. In this paper we present modifications to MRPerf, an existing tool used to simulate MapReduce in data center clouds, enabling it to simulate Hadoop MapReduce jobs in edge networks. Our results indicate that the default Hadoop scheduler does not perform optimally in an edge environment and should be replaced. We also conclude that bandwidth is not a limiting factor in some MapReduce jobs for capacities that are even available in some residential areas. Aside from scheduling, our experiments show that latency for MapReduce jobs is a large factor of performance that is challenging to improve upon. The reason for this is that data is always node remote due to the nature of an edge-based system.

1 Introduction

Edge based cloud computing is a combination of two popular ideas, edge computing and cloud computing. We believe that the combination of these two ideas offers the potential for a high performance per dollar ratio by leveraging the enormous amount of unused computational power in idle home computer processors that have access to always-on internet connections. Despite its potential this type of computing paradigm is not being used and is only sparsely investigated. This paper presents simulation results for MapReduce [2] jobs in an edge based cloud. The high level of coordination and communication involved in MapReduce extends beyond the usual peer-to-peer and grid computing type of computation associated

with edge computing. For our simulation environment we investigated MRPerf, a MapReduce simulation tool built on top of the network simulator ns-2, which was developed by researchers at Virginia Tech and IBM Almaden Research Center. To simulate an edge based network we modified parameters in our configuration to better mimic a peer-to-peer type system of nodes without central infrastructure. Each node has fewer resources and is located at a further distance from one another than nodes in a data center.

This paper has the following structure: Section 2 discusses some related work, Section 3 presents an overview of the tools used for our simulations, Section 4 discusses the parameters that affect MapReduce jobs, Section 5 discusses our implementation of the edge based simulator, Section 6 discusses our simulations and results, Section 7 presents the conclusions drawn from the simulation results, and Section 8 discusses potential future research in the area of edge based cloud computing.

2 Related Work

Zaharia et al.[8] presented LATE, a scheduling algorithm to handle heterogeneity in a data center environment. LATE attempts to schedule tasks based on the longest approximate time to completion. The approximation is based on the heterogeneity of the nodes and the current progress of the task. The work presented in [8] could potentially be extended to schedule tasks in an edge based cloud where it includes latency and node bandwidth in its approximate time to completion calculations.

In general, work applying cloud computing concepts, such as those that drive the MapReduce pro-

gramming model, directly to edge networks is sparse and this area of research is still in a stage of early development. Because of this our work focused more on the feasibility of edge based cloud computing, MapReduce in particular, from the aspect of performance and in doing so we ignored many practical issues such as security and privacy.

3 Tools

We leveraged two existing projects for our work. The first is MRPerf [6, 7]. It was designed to simulate MapReduce jobs when provided with a network topology for a data center. MRPerf can measure the performance of these MapReduce jobs with a good amount of accuracy and deliver performance information to the user.

We considered the use of a topology generator such as GT-ITM [3] or BRITE [1] to provide us with a WAN model but decided that we would obtain more realistic results if we used information about the actual Internet instead of generating router topologies. We decided to use the data collected from the Rocketfuel [9] project which mapped a significant portion of the Internet. These data provided us with a realistic topology for edge based simulations.

3.1 MRPerf

We chose to use MRPerf because of its impressive implementation of MapReduce and the accuracy of its simulation results. Our goal was to evaluate an edge computing network with computationally demanding, highly coordinated applications. MapReduce is a good tool for testing this because its design requires participating machines to break up a computationally intensive job into separate tasks, schedule the tasks efficiently, process the tasks, and then merge them back together upon completion. MRPerf merges MapReduce and network simulation to achieve a seamless simulation environment. It accomplishes this by utilizing the popular packet level network simulator ns-2 [4] and the Hadoop [5] scheduling algorithm to schedule tasks. MRPerf claims to have the ability to simulate MapReduce operations in a wide range of data center topologies. For data center topologies MRPerf claims to predict

simulation performance within 5.22% of actual measurements for map and 12.83% for reduce for a double rack cluster with 16 to 128 cores [6]. Although MRPerf was designed to simulate MapReduce jobs for multi-rack data centers with many nodes per rack and high bandwidth single hop links connecting them together it allows some flexibility since it was built on top of the ns-2 network simulator. ns-2 has already proven to be capable of simulating a large variety of networks including networks based on grid computing and peer to peer type topologies [13, 3]. In our work we take advantage of this by creating a topology that closely resembles the internet and passing that to MRPerf’s topology reader (Figure 1).

3.2 Rocketfuel

We wanted to run our simulation on a fairly large network to get an accurate idea of performance when route scheduling and link latency are taken into consideration. Rocketfuel was a project completed in 2003 at the University of Washington that mapped ISP topologies using traceroutes from 800 different vantage points. They mapped over 50 thousand IP addresses which represented about 45 thousand routers in 537 POPs connected by 80 thousand links. Their maps cover ISPs in the United States, Europe, Australia, and India. They claimed that other than the maps made by the ISPs themselves the maps produced by Rocketfuel were the most detailed maps available in 2003. We used the link latencies discovered by the Rocketfuel project to as input to our simulations. This allows us to use a more accurate representation of a real edge based network for our simulations.

4 Simulation Design

In addition to using an accurate edge topology we also needed to modify some of the parameters used in MRPerf to simulate MapReduce tasks. MRPerf simulates MapReduce tasks based on three classes of parameters: cluster parameters, configuration parameters, and framework parameters. The rest of this section is dedicated to describing these parameters as well as modifications that could be made to improve performance in edge networks. Section 5

describes the modifications that we actually made to MRPerf for our evaluations.

4.1 Cluster parameters

MRPerf is designed to work with a data center using racks which consist of many nodes that are close to each other in the network, this means that many times packets are one hop away from their destination. These racks of nodes are connected together to form a cluster. In an edge computing network there is no concept of a centralized cluster of computational power. Edge computing networks use PCs that are connected to the Internet via a LAN through a router or gateway. This means that a packet's destination is almost always several network hops away. The racks that make up a data center cluster have many compute nodes each of which have several processors whereas a PC in an edge network has anywhere from one to four processors. Another difference in the cluster parameters is related to node heterogeneity. A data center setup usually has homogeneous nodes, whereas in an edge network nodes have a high level of heterogeneity. The interconnect topologies also play a role in the cluster parameters. The inter/intra-rack topology of a data center is much less complex than the Internet.

4.2 Configuration parameters

Configuration parameters are related to the MapReduce jobs themselves. They adjust settings such as data chunk size, data replication, tasks per node, and tasks in each job. The default Hadoop settings are optimized for data centers. Some of the changes to the configuration parameters depend upon the optimal configuration of the edge computing network. Whether data chunk size is large or small depends upon whether the edge computing system performs better when I/O is optimized or when parallelism is increased. The cost of data replication is clearly higher in an edge computing network than in a data center so the cost of redundancy should be weighed against the cost of rescheduling failed tasks.

4.3 Framework parameters

Framework parameters deal with data placement and scheduling algorithms. The distance of nodes relative

the machine requesting jobs needs to be considered when partitioning data across available resources. In a data center with many compute nodes the wrong choice of data placement can affect performance to a degree, but is not nearly as costly as the wrong choice when compute nodes could potentially be in a different country. Wang et al.[6] compares three types of data locality: node-local, rack-local, and rack-remote. Our simulations used the rack-remote algorithm in MRPerf, which was the slowest of all of the data center evaluations [6].

The scheduling algorithm can also be optimized for edge networks to account for node heterogeneity, bandwidth capacity, and round trip time. As discussed above the proper balance of data chunk size, data placement, and task scheduling are important so that the system does not suffer huge performance losses when transferring large amounts of data over relatively slow links to computationally slow nodes.

4.4 Limitations

Some of the inherent limitations of MRPerf are worth discussing. These limitations are related to disk storage and simulation of node reliability. The former works to our advantage in this case. MRPerf only allows for a single storage device per node. In a data center this might be an unusual constraint however since the average PC that constitutes our edge computing network will have only a single storage disk this should not present a problem in our evaluation. The early release of MRPerf that is publicly available does not currently support realistic node failures or lagging nodes. Ideally, we would have been able to simulate lagging nodes and failed nodes in our evaluation. However, without an accurate mechanism to simulate these cases we leave this for future work.

5 Implementation

The goal of our implementation was to discover if the MapReduce framework is feasible in an edge network. The implementation used in our simulations addressed some but not all of the issues discussed in Section 4. We discuss the issues that have not been addressed in Section 8.

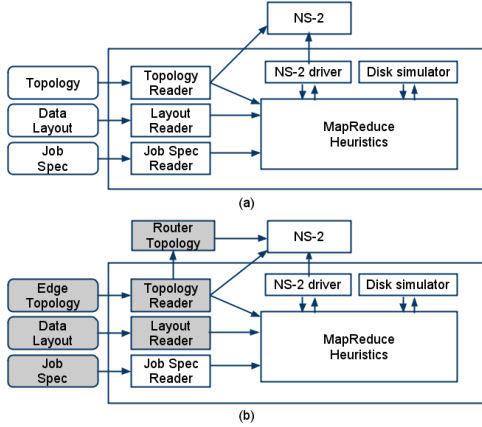


Figure 1: (a) shows the original MRPerf architecture. The shaded components in (b) were changed or added for the edge based MapReduce simulations.

5.1 MRPerf Input and Pre-Processing Overview

Figure 1 (a) shows the original MRPerf architecture. The topology, data layout, and job specifications are provided as input in the form of XML and TCL code. The topology is provided in XML format which is parsed by a Python script. This script outputs the necessary TCL code for use in ns-2 and an XML data layout file which is read directly by the MRPerf module integrated into ns-2. In addition to the XML input which specifies the topology the TCL input contains the MapReduce job specification such as the number of map and reduce slots available for each node. These variables are used directly by the simulator.

Figure 1 (b) highlights the changes and additions to the MRPerf architecture that were required to simulate the edge based networks. The basic idea behind our modifications was to create "racks" each containing one node. These one node "racks" could then use the remote rack scheduling algorithm to receive data and tasks across the network. The following discussion outlines the modifications made to each part of the MRPerf architecture in order to realize this idea.

Topology and Topology Reader. The original construction of MRPerf accounted for the CPU speed, amount of memory, and I/O speed but forced all nodes participating in a computation to be homo-

CPU speed (GHz)	1.5, 1.6, 1.8, 2.0, 2.3, 2.4, 2.5, 3.0, 3.2
Number of CPUs	1, 2
Number of cores per CPU	1, 2, 4
Number of disks	1, 2
Disk capacity (GB)	40, 60, 80, 100, 120, 160, 180, 200, 250, 300, 400, 450, 500, 550, 600, 650, 700, 750, 800, 850, 900, 950, 1000
Disk read bandwidth (MB/s)	250, 260, 270, 280
Disk write bandwidth (MB/s)	60, 65, 70, 75
NIC capacity	100Mbps, 1Gbps
Memory capacity (GB)	0.5, 1, 2, 3, 4, 5, 6

Table 1: Summary of ranges used in host specification

geneous. We changed the semantics of the topology reader to allow a different host configuration for each rack. We then employed the use of a script to generate sets of end hosts which we attached to the routers. Table 1 summarizes the range of specifications that we used in our evaluation. In addition to including a range of host specifications we modified the Topology Reader to randomly choose the bandwidth of the last hop from the Internet router to the host based on a normal distribution. In our evaluation we set the mean of the distribution at 16 Mbps with a standard deviation of 8 Mbps.

Data Layout and Layout Reader. The data layout components of MRPerf control size of data chunks and the initial scheduling of data chunks to individual disks. The semantic modification we made in the topology portion of MRPerf forced us to also modify the initial chunk distribution.

Router Topology. The original Topology Reader component of MRPerf generated a complete data center topology script for use in ns-2 based on a specification provided in the form of an XML doc-

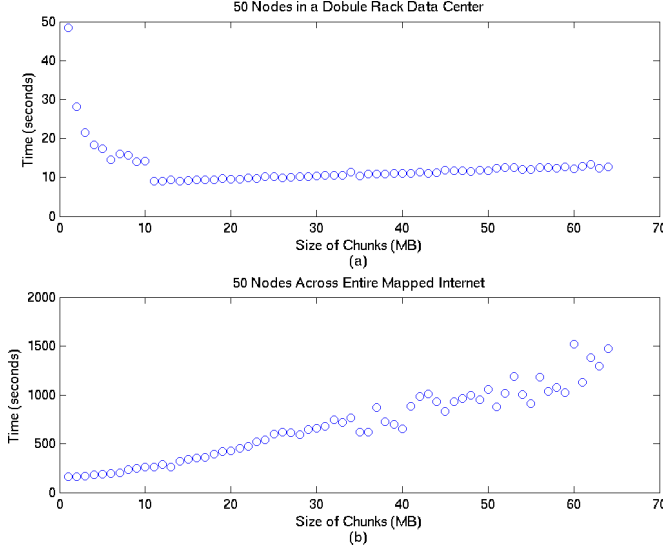


Figure 2: Results for 50 node topologies in a double rack data center configuration (a) and a random configuration distributed across the entire mapped portion of the Internet (b) in which the size of the data chunks was variable.

ument. For our simulations we directly converted the topologies provided by the Rocketfuel results into TCL code and simply made the Topology Reader output a call to this code.

6 Simulations and Results

We conducted simulations in which we varied several aspects of the simulation including the number of nodes, chunk size, link bandwidth, and the number of map and reduce slots per node. The results presented in this section are the average over five simulations given the same parameters the only variable being the specific property that was being tested. All of the hosts in the edge simulations have specifications taken at random from the specifications presented in Table 1. Each simulation consisted of a MapReduce job sorting one gigabyte of data. All statistically insignificant data points have been removed.

For the chunk size simulations we conducted two types of simulations, one in which the nodes are in a double rack data center and one in which nodes are randomly distributed throughout the mapped portion of the Internet. For both types we used 50 node topologies which we held constant for all simulations.

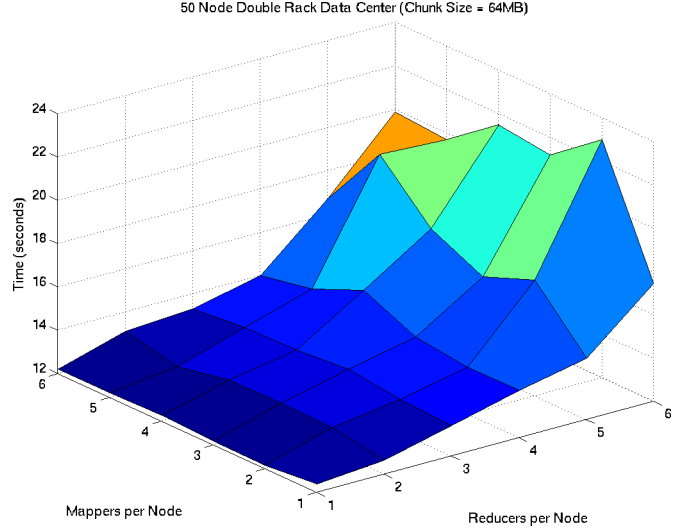


Figure 3: Results for 50 node double rack data center topology in which the number of map and reduce slots available on each host was variable.

The total MapReduce run times calculated for these simulations are presented in Figure 2.

For the simulations in which we varied the number of available map slots and the available number of reduce slots per node we used a chunk size of 6 MB and a 50 node topology randomly distributed across the mapped portion of the Internet. The topology was held constant as we varied the number of available map and reduce slots per node. The results from these simulations are presented in Figure 4. We also ran similar simulations on a 50 node double rack data center configuration with a 64 MB chunk size. These results are presented in Figure 3.

For the simulations in which we varied the bandwidth of the link connecting the end hosts to the routers we used a 50 node topology where nodes were distributed across the entire mapped Internet. This topology was held constant as we varied the mean of the normal distribution that was used to assign the bandwidth of the end hosts' links. The standard deviation of the normal distribution was held at 8 Mbps. The results from these simulations are presented in Figure 5.

We ran simulations that varied the number of nodes in a network for three different types of topologies. We varied the number of nodes that were distributed across the mapped portion of the Internet, the number of nodes that were distributed across one

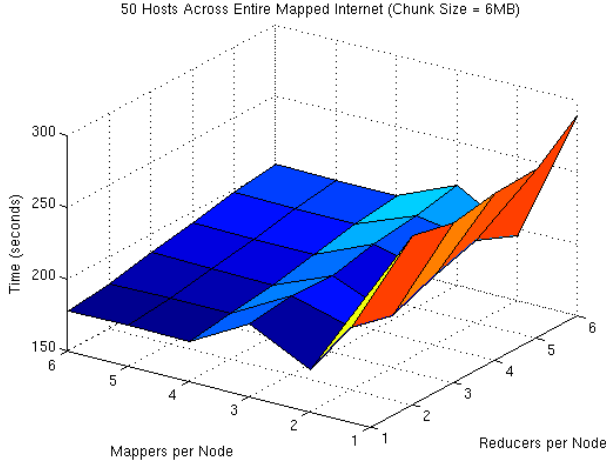


Figure 4: Results for a 50 node topology randomly distributed across the mapped portion of the Internet in which the number of map and reduce slots available on each host was variable.

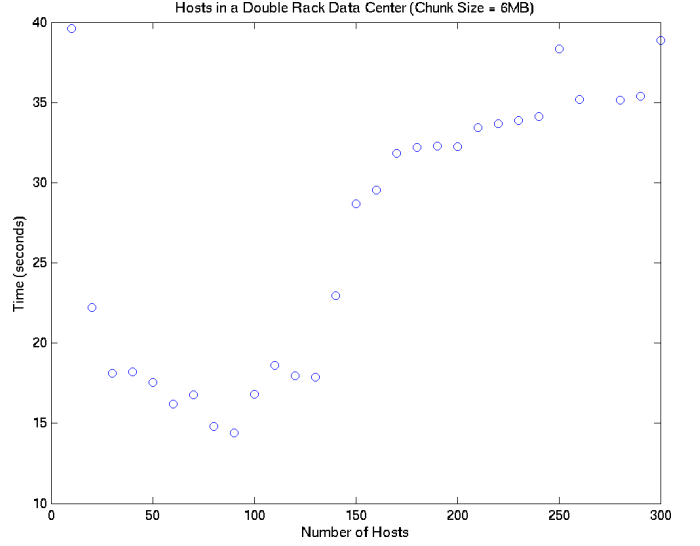


Figure 6: Results for double rack data center topologies in which the total number of nodes was variable.

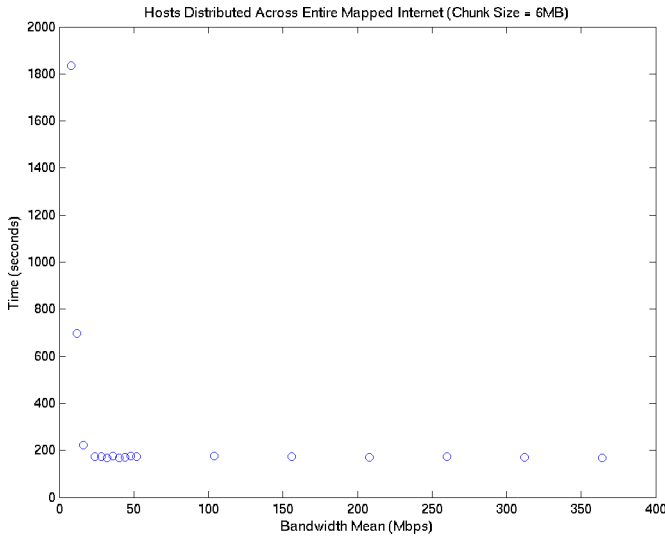


Figure 5: Results for a 50 node topology distributed across the mapped portion of the Internet in which the mean of the normal distribution that we used to randomly assign the bandwidth of the link connected to each node varies.

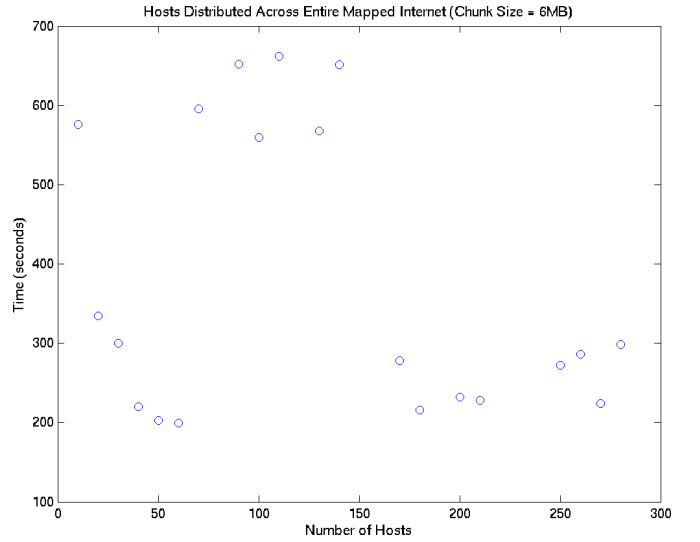


Figure 7: Results for topologies in which each consecutive point represents the running time when 10 additional hosts are randomly added to the previous topology. Hosts are distributed across the entire mapped portion of the Internet.




Figure 8: Results for topologies in which each consecutive point represents the running time when 10 additional hosts are randomly added to the previous topology. Hosts are distributed across AS 11537 which spans the entire United States and parts of southern Canada.

autonomous system consisting of 15 mapped routers which span the United States, and the number of nodes in a double rack data center. Each consecutive simulation for the edge based data kept the previous simulation’s topology and added 10 additional nodes randomly distributed throughout the topology. The results for the data center and edge based simulations are presented in Figures 6, 7, and 8.

7 Conclusions

From our results we can conclude that the data chunk size used in edge based MapReduce jobs should be considerably smaller than what is usually used in data center MapReduce jobs. Figure 2 shows that while certain data center topologies can suffer if the chunk size is too small the edge based network produces the best results when the chunk size is small.

The results from our simulations varying the available number of map and reduce slots on each node (Figure 4) show that edge-based MapReduce jobs benefit when the number of available map slots on each node increases but does not necessarily benefit as the number of reduce slots available on each

node increases. Figure 3 indicates that increasing the number of map slots available on each node does not affect the performance in the double rack data center topology but increasing the number of reduce slots available has a negative impact on the performance. This suggests that the number of map slots and reduce slots should be configured carefully not only to suit the type of MapReduce job that is running but also based on the configuration of the topology.

We can conclude from our bandwidth simulations (Figure 5) that adding additional bandwidth to the end host links is no longer beneficial to the MapReduce job performance after a certain bandwidth is reached. For the topology used in the simulation this point is at about 20 Mbps. This suggests that bandwidth could be a limiting factor for the other simulations that we ran which used a mean host link bandwidth of 16 Mbps.

By varying the number of hosts in a double rack data center (Figure 6) we can see that the total run time increases as more nodes are added after a certain point for the data center topology evaluated. This is probably caused because the sort algorithm becomes I/O bound and no longer benefits from extra computational power. It is more difficult to draw conclusions about the effects of adding more nodes in an edge based cloud (Figures 7 and 8). Perhaps the only definite conclusion that we can draw is that the default scheduler should be replaced by one that is aware of the network topology and node specifications. In our evaluation it was possible for the scheduler to do very well with one topology and very poorly when a distant node is added which the scheduler blindly scheduled tasks to. The development of this new scheduler is left for future work.

8 Future Research

We did not present results related to modifications to the Hadoop scheduling algorithm used by MRPerf. This is primarily because the early release of MRPerf we used for our simulations does not support realistic node failures. We are confident that a scheduling algorithm that is aware of the round trip times to each node and heterogeneity of the nodes would perform better than the default Hadoop scheduling algorithm. An analysis that involves the simulation

of node churn is also left for future work.

A possible alternate simulation tool to MRPerf is CloudSim [10]. Future research could investigate the use of MapReduce with this Java based simulation tool that has recently been developed. CloudSim was built on top of the existing tools, JavaSim and GridSim, which have already proven themselves valid for simulation. CloudSim itself is an effort to add a new layer of features on top of these tools that are more critical and immediately relevant to cloud computing. By abstracting away all the lower levels CloudSim can focus on doing its job correctly and at the same time trust that the layers below are functioning correctly as well. The results gathered from simulation with this tool would be a valuable comparison to the results we present in this paper. Other avenues for research could address the security and privacy implications for this type of computing paradigm.

References

- [1] ALBERTO MEDINA, ANUKOOL LAKHINA, I. M., AND BYERS, J. Brite: Universal topology generation from a user's perspective. (user manual), Apr 2001.
- [2] DEAN, J., AND GHEMAWAT, S. Mapreduce: Simplified data processing on large clusters. In *OSDI'04: Sixth Symposium on Operating System Design and Implementation* (Dec 2004).
- [3] ELLEN W. ZEGURA, K. C., AND BHATTACHARJEE, S. How to model an internetwork. In *Proceedings of IEEE Infocom '96* (1996).
- [4] FALL, K., AND VARADHAN, K. The network simulator - ns-2. ns manual.
- [5] FOUNDATION, A. S. Hadoop documentation, Apr 2009.
- [6] GUANYING WANG, ALI R. BUTT, P. P., AND GUPTA, K. A simulation approach to evaluating design decisions in mapreduce setups. In *Proceedings of the 17th Annual Meeting of the IEEE/ACM International Symposium on Modelling, Analysis and Simulation of Computer and Telecommunication Systems* (Sep 2009).
- [7] GUANYING WANG, ALI R. BUTT, P. P., AND GUPTA, K. Using realistic simulation for performance analysis of mapreduce setups. In *Proceedings of the 1st ACM workshop on Large-Scale System and Application Performance* (Jun 2009).
- [8] M. ZAHARIA, A. KONWINSKI, A. J. R. K., AND STOICA, I. Improving mapreduce performance in heterogeneous environments. In *8th USENIX Symposium on Operating Systems Design and Implementation* (Dec 2008).
- [9] NEIL SPRING, R. M., AND ANDERSON, T. Measuring isp topologies with rocketfuel. In *SIGCOMM '02* (Aug 2002).
- [10] RAJKUMAR BUYYA, R. R., AND CALHEIROS, R. N. Modeling and simulation of scalable cloud computing environments and the cloudsims toolkit: Challenges and opportunities. In *Proceedings of the 7th High Performance Computing and Simulation Conference* (Jun 2009).
- [11] ROSE, C. A. F. D., AND BUYYA, R. Cloudsim: A novel framework for modeling and simulation of cloud computing infrastructures and services. Tech. rep., Grid Computing and Distributed Systems Laboratory, The University of Melbourne, Australia, Mar 2009.
- [12] WANG, G., AND BUTT, A. R. Mrperf simulator. mrperf readme.
- [13] WU KUN, DAI HAN, Y. Z. S. L. Ndp2psim: A ns2-based platform for peer-to-peer network simulations. In *Proc. of International Workshop on Service Grid Computing and Applications* (Nov 2005).